


Problem 1.


- Find the equation of the hyperplane (in terms of w) WITHOUT solving a quadratic programming problem. Make a sketch of the problem.
- Calculate the margin.
- Find the α 's of the SVM for classification.
- Perform SVM with Matlab or R with RBF kernel and plot the result. Is this result different from the one you obtained by hand?

Solution

Part (a)



(a) Original Data



(b) Transformed Data

From this transformed data, we can clearly see our two support vectors (one for each class). For the orange class (-1) this is at point $[1, 1]$ and for the blue class (+1) this is at point $[2, 2]$. SVM requires that all points be on their respective "positive" sides of the hyperplane, thus we can see that this hyperplane must pass between the two support vectors we have chosen. This means, that if we imagine a line segment connecting the points $\{[1, 1], [2, 2]\}$ our hyperplane MUST intersect that line segment. Further, as these points are support vectors, they must lie at equal distance to the hyperplane, from these observations we conclude that the hyperplane will intersect these two points at their midpoint, this will be the point $\frac{1}{2}[1 + 2, 1 + 2] = [\frac{3}{2}, \frac{3}{2}]$.

As our transformed data exists in two dimensions, our hyperplane will take the form of a line, as we know an intercept of this line, the only thing that remains to be found is the

slope of the line, once we have that we can derive an equation for this line in the desired form.

Intuitively we would expect the hyperplane to be perpendicular to our imaginary line segment connecting our two support vectors. To see why this must be the case, simply consider the case where it is not. At one extreme, our hyperplane is in line with both support vectors, producing a margin equal to 0, as we rotate our line, this margin will increase in size. It will continue to increase in size until it reaches the perpendicular we are considering, if we continue to rotate beyond the perpendicular, each edge will begin to approach the opposite class, thus resulting in a decrease again in our margin. Thus maximum margin is achieved at the perpendicular. As the line segment through $\{[1,1],[2,2]\}$ has slope 1, this implies our hyperplane has slope -1.

In standard form this means our line will take the form:


$$ax + b = -1x + 3$$

Some simple algebra will produce the vectorized form:

$$\begin{aligned} -x + 3 &= y \\ x + y - 3 &= 0 \\ [1; 1]^T [x; y] - 3 &= 0 \end{aligned}$$

From this we have $w = [1, 1]$ and $b = -3$.

To check our work we need to consider our class labels (else we would simply multiply both sides by -1), we expect $w'[2; 2] + b = 1$ and $w'[1; 1] + b = -1$. Plugging these values into matlab produces the expected results so I'll assume I did my work correctly.



Part (b)

The margin will simply be the distance between my two support vectors.

$$\begin{aligned} \text{dist}([1, 1], [2, 2]) &= \text{sqr}t((2 - 1)^2 + (2 - 1)^2) \\ &= \text{sqr}t(2(1)^2) \\ &= \text{sqr}t(2) \end{aligned}$$

Thus we find we have a margin of $\text{sqr}t(2)$.

Part (c)

Note that the majority of our points are not support vectors, for all of these points we have that $\alpha = 0$. Note that our support vectors are observation 1 and 5, thus we still need to find α_1 and α_5 . From my notes we have that $w^\top = \sum_{i=1}^m \alpha_i Y_i x_i^\top$, noting that each class only has one support vector we can simplify this as follows

$$w^\top = \alpha_1 * [2; 2]^\top + -\alpha_5 * [1; 1]^\top$$

This is a 2x2 linear system, format it as such.

$$\begin{bmatrix} 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 2 & -1 \\ 2 & -1 \end{bmatrix} \begin{bmatrix} \alpha_1 \\ \alpha_5 \end{bmatrix}$$

From here it is pretty obvious that $\alpha_5 = \alpha_1$ denote this δ and we have $1 = 2\delta - 1\delta = (2 - 1)\delta = \delta$. Thus $\delta = 1$.

$$\begin{aligned} \alpha_1 &= 1 \\ \alpha_5 &= 1 \\ \alpha_i &= 0 \text{ otherwise} \end{aligned}$$

Part (d)

```
function p1
```

```
clear all;
close all;
hold on;
```

```
%% Data
```

```

data = [
    % Class 1
    2, 2, 1;
    2, -2, 1;
    -2, -2, 1;
    -2, 2, 1;
    % Class -1
    1, 1, -1;
    1, -1, -1;
    -1, -1, -1;
    -1, 1, -1
];

%% Manual SVM

axis([-2 10 -2 10]);

classA = data(data(:,3)==1, 1:2);
classB = data(data(:,3)==-1, 1:2);

classA = phi(classA);
classB = phi(classB);

scatter(classA(:,1), classA(:,2), 'LineWidth', 2)
scatter(classB(:,1), classB(:,2), 'LineWidth', 2)

% Plot the decision boundary.
x = [-2, 10];
y = -1 * x + 3;
[1; 1]' * x' - 3
plot(x, y, 'k', 'LineWidth', 2);

% Plot the margins.
y1 = -1 * x + 3 + 1;
y2 = -1 * x + 3 - 1;
plot(x, y1, '--k', 'LineWidth', 1);
plot(x, y2, '--k', 'LineWidth', 1);


%% Matlab SVM

% X = data(:,1:2);
% Y = data(:,3);
%
% mdl = fitcsvm(X, Y, 'KernelFunction', 'rbf', 'OptimizeHyperparameters', 'auto',...
```

```
%      'HyperparameterOptimizationOptions', struct('Optimizer', 'bayesopt', 'Kfold', 10  
% plotsum mdl)  
  
end
```

Problem 2.

Verify by means of simulation that each bootstrap sample will contain $1 - 1/e \approx 63.2$ of the original sample.

Solution

```
function p2

clear all;
close all;
hold on;

M = 10000; % How many simulations are we going to run?
ans = zeros(M,1); % This should approach 1 - 1/e.
actual = (1 - 1/e); % The actual result we are expecting.

% Run M bootstrap samples.
for n = 1:M
    % Random normal data set of size 'n'.
    X = randn(n, 1);

    % Generate random indecies in the range [1,n].
    idx = ceil(rand(n,1)*n);

    % The number of samples contained in X is equal
    % to the number of UNIQUE indecies in idx.
    % To get the fraction simply divide by n.
    ans(n) = length(unique(idx)) / n;
end


plot(1:M, ans, 'LineWidth', 2);
plot([1 M], [actual actual], 'LineWidth', 2);

end
```

Problem 3.

Compare SVM and LS-SVM on the ozon level detection data set on Blackboard. This dataset contains 72 measurement variables and was measured between 1/1/1998 and 12/31/2004. All missing values have been removed. The goal is to detect whether there was too much ozon (class label 0) or a normal day (class label 1). The class label is the last variable. Set up the simulation and clearly describe what you are doing and why. Finally, state, according to your findings (boxplots, ROC curves, etc.), the best classifier for this problem.

Solution



```
function p3
```

```
clear all;
close all;
hold on
```

```
%% Read the data from the given spreadsheet.
data = xlsread('ozon.xlsx');
```

```
X = data(:,1:72); % All of our data, but no labels (1874 Entries).
Y = data(:,73);    % Labels corresponding to each entry in X.
```

```
%% SVM
```



```

tic;

% Kernel Functions: rbf, linear, or polynomial
mdl = fitcsvm(X, Y, 'KernelFunction', 'rbf', 'OptimizeHyperparameters', 'auto',...
    'HyperparameterOptimizationOptions', struct('Optimizer', 'bayesopt', 'Kfold', 10));
% linear and polynomial kernels weren't working with the extra parameters
% but it worked after removing them
% mdl = fitcsvm(X, Y, 'KernelFunction', 'linear');

% Get our misclassification rate estimate through cross validation.
kfoldLoss(crossval(mdl))
toc;

% rbf                0.0671 # 148.50 seconds
% linear            0.0693 # 334.09 seconds
% polynomial        NaN      # 2.0702 seconds

%% LSSVM

tic;

% Kernel Functions: lin_kernel, poly_kernel, RBF_kernel
mdl = initlssvm(X, Y, 'c', [], [], 'poly_kernel', 'p');
mdl = tunelssvm(mdl, 'simplex', 'crossvalidatelssvm', {10, 'misclass'});
mdl = trainlssvm(mdl);
% Get our misclassification rate estimate through cross validation.
crossvalidate(mdl, 10, 'misclass')

toc;

% lin_kernel        0.0693 # 48.76 seconds
% poly_kernel       0.1012 # 43.01 seconds
% RBF_kernel        0.0596 # 75.90 seconds

%% Box Plots

B = 100;
mcrcsvm = zeros(B,1);
mcrlssvm = zeros(B,1);

tic;

% Basically just copy and paste the prof's example code.
for b = 1:B

```

```

% For sake of performance, use a random sample size
% of 500 (375 Train, 125 Test).
% We already have an estimate of misclassification error,
% so we really are only interested in the variance, so this should be fine.

N = 500; % How big is the random subset we are considering?
r = randperm(size(X,1)); % Need a random permutation over ALL indecies...
r = r(1:N); % but only want N when we are all said and done.
ntr = ceil(0.75 * N); % The first 75% are for training.

% Break up our data into training...
Xtr = X(r(1:ntr), :);
Ytr = Y(r(1:ntr));

% and test sets.
Xtest = X(r(ntr+1:N), :);
Ytest = Y(r(ntr+1:N));

%%%%%%%%%
% SVM %
%%%%%%%%%
mdl = fitcsvm(X, Y, 'KernelFunction', 'rbf', 'OptimizeHyperparameters', 'auto',...
    'HyperparameterOptimizationOptions', struct('Optimizer', 'bayesopt', 'Kfold', 10));
yh = predict(mdl, Xtest);
mcrcsvm(b, 1) = sum(Ytest ~= yh) / size(Ytest, 1);

%%%%%%%%%
% LSSVM %
%%%%%%%%%
mdl = initlssvm(X, Y, 'c', [], [], 'RBF_kernel', 'p');
mdl = tunelssvm(mdl, 'simplex', 'crossvalidatelssvm', {10, 'misclass'});
mdl = trainlssvm(mdl);

% For whatever reason lssvm is only letting me do one prediction at a time,
% so take this pain in the ass approach to get our predictions.
T = size(Xtest, 1);
yh = zeros(T, 1);
for idx = 1:T
    yh(idx, 1) = predict(mdl, Xtest(idx, :), 1);
end

mcrlssvm(b, 1) = sum(Ytest ~= yh) / size(Ytest, 1);

close all;

```

```
end

toc;

legend = { 'SVM', 'LSSVM' };
boxplot([mcsvm, mcrlssvm], legend);
ylabel('Misclassification rate on test data');

end
```
